

```
#required libraries

import numpy as np
import math
import matplotlib.pyplot as plt
import matplotlib.colors
import time

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, mean_squared_error, log_loss
from tqdm import tqdm_notebook

from IPython.display import HTML
import warnings
from sklearn.preprocessing import OneHotEncoder
from sklearn.datasets import make_blobs

import torch
warnings.filterwarnings('ignore')

torch.manual_seed(0)

<torch._C.Generator at 0x7f1bce072090>


my_cmap = matplotlib.colors.LinearSegmentedColormap.from_list("", ["red" , "yellow" , "green"])
```

▼ Generate dataset

```
data , labels = make_blobs(n_samples = 1000 , centers = 4 , n_features= 2 , random_state= 0 )  
print(data.shape , labels.shape)
```

```
(1000, 2) (1000,)
```

```
plt.scatter(data[:, 0] , data[:, 1] , c = labels , cmap = my_cmap )  
plt.show()
```



```
X_train , X_val , Y_train , Y_val = train_test_split(data , labels , stratify = labels , random_state = 0)
```

```
print(X_train.shape , Y_train.shape , labels.shape )
```

```
(750, 2) (750,) (1000,)
```

▼ Using torch tensors and autograd

```
X_train , Y_train , X_val , Y_val = map(torch.tensor , (X_train , Y_train , X_val , Y_val ))
print(X_train.shape , Y_train.shape)
```

```
torch.Size([750, 2]) torch.Size([750])
```

```
def model(x):
    a1 = torch.matmul(x, weights1) + bias1 # (N, 2) x (2, 2) -> (N, 2)
    h1 = a1.sigmoid() # (N, 2)
    a2 = torch.matmul(h1, weights2) + bias2 # (N, 2) x (2, 4) -> (N, 4)
    h2 = a2.exp()/a2.exp().sum(-1).unsqueeze(-1) # (N, 4)
    return h2
```

```
y_hat = torch.tensor([[0.1 , 0.2, 0.3 , 0.4] , [0.8 , 0.1 , 0.05 , 0.05]])
y = torch.tensor([2, 0])
(-y_hat[range(y_hat.shape[0]) , y ].log()).mean().item()
(torch.argmax(y_hat ,dim =1 ) == y).float().mean().item()
```

```
0.5
```

```
def loss_fn(y_hat, y):
    return -(y_hat[range(y.shape[0])], y).log()).mean()
```

```
def accuracy(y_hat , y ) :
    pred = torch.argmax(y_hat , dim =1)
    return (pred == y ).float().mean()
```

```
plt.style.use("seaborn")
```

```
torch.manual_seed(0)

#initialize the weights and biases using He Initialization
weights1 = torch.randn(2, 2) / math.sqrt(2)
weights1.requires_grad_()
bias1 = torch.zeros(2, requires_grad=True)

weights2 = torch.randn(2, 4) / math.sqrt(2)
weights2.requires_grad_()
bias2 = torch.zeros(4, requires_grad=True)

#set the parameters for training the model
learning_rate = 0.2
epochs = 10000

X_train = X_train.float()
Y_train = Y_train.long()
X_val = X_val.float()
Y_val = Y_val.long()

loss_arr = []
acc_arr = []
val_acc_arr = []

#training the network
for epoch in range(epochs):
    y_hat = model(X_train) #compute the predicted distribution
    loss = loss_fn(y_hat, Y_train) #compute the loss of the network
    loss.backward() #backpropagate the gradients
    loss_arr.append(loss.item())
    acc_arr.append(accuracy(y_hat, Y_train))

    with torch.no_grad(): #update the weights and biases
        val_acc_arr.append(accuracy(model(X_val), Y_val))

        weights1 -= weights1.grad * learning_rate
        bias1 -= bias1.grad * learning_rate
        weights2 -= weights2.grad * learning_rate
```

```
bias2 -= bias2.grad * learning_rate
weights1.grad.zero_()
bias1.grad.zero_()
weights2.grad.zero_()
bias2.grad.zero_()

plt.plot(loss_arr, 'r-', label='loss')
plt.plot(acc_arr, 'b-', label='train accuracy')
plt.plot(val_acc_arr, 'g-', label='val accuracy')
plt.title("Loss plot - Using tensors and autograd")
plt.xlabel("Epoch")
plt.legend(loc='best')
plt.show()
print('Loss before training', loss_arr[0])
print('Loss after training', loss_arr[-1])
```

▼ Using nn Functional

```
import torch.nn.functional as F
```

```
10
```

```
torch.manual_seed(0)
weights1 = torch.randn(2,2) / math.sqrt(2)
weights1.requires_grad_()
bias1 = torch.zeros(2, requires_grad = True)
```

```
weights2 = torch.randn(2,4) / math.sqrt(2)
weights2.requires_grad_()
bias2 = torch.zeros(4 , requires_grad = True)
```

```
learning_rate = 0.2
epochs = 10000
```

```
loss_arr = []
acc_arr = []
```

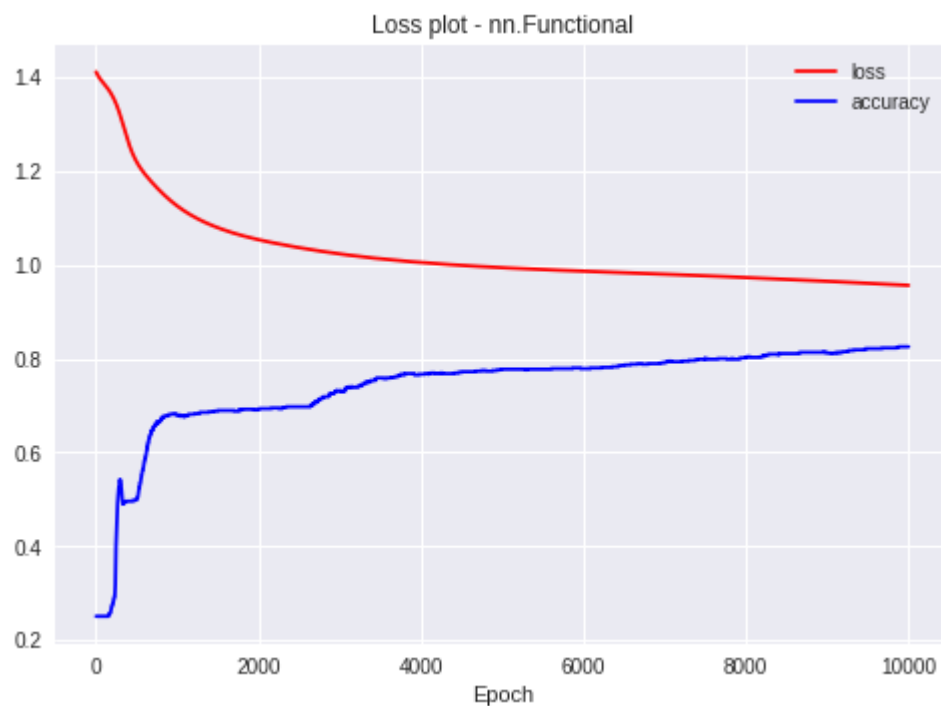
```
for epoch in range(epochs) :
    y_hat = model(X_train)
    loss = F.cross_entropy(y_hat , Y_train)
    loss.backward()
    loss_arr.append(loss.item())
    acc_arr.append(accuracy(y_hat , Y_train))
```

```
with torch.no_grad():
    weights1 -= learning_rate*weights1.grad
    bias1 -= bias1.grad*learning_rate
    weights2 -= learning_rate*weights2.grad
    bias2 -= bias2.grad*learning_rate
```

```
weights1.grad.zero_()
bias1.grad.zero_()
weights2.grad.zero_()
bias2.grad.zero_()
```

```
plt.plot(loss_arr , "r-" , label = "loss")
plt.plot(acc_arr , "b-" , label = "accuracy")
plt.legend(loc = "best")
plt.title("Loss plot - nn.Functional")
plt.xlabel("Epoch")
plt.show()
```

```
plt.show()
print('Loss before training', loss_arr[0])
print('Loss after training', loss_arr[-1])
```



```
Loss before training 1.411980199813843
Loss after training 0.9561843276023865
```

▼ Using NN Parameters

```
import torch.nn as nn
```

```
class FirstNetwork(nn.Module):
    def __init__(self) :
        super().__init__()
        torch.manual_seed(0)
        self.weights1 = nn.Parameter(torch.randn(2,2) / math.sqrt(2))
        self.bias1 = nn.Parameter(torch.zeros(2))
        self.weights2 = nn.Parameter(torch.randn(2,4) / math.sqrt(2))
        self.bias2 = nn.Parameter(torch.zeros(4))
```

```
def forward(self ,X):
    a1 = torch.matmul(X , self.weights1 ) + self.bias1
    h1 =a1.sigmoid()
    a2 = torch.matmul(h1 , self.weights2) + self.bias2
    h2 = a2.exp()/a2.exp().sum(-1).unsqueeze(-1)
    return h2
```

```
def fit(epochs = 10000 , learning_rate = 0.2 , title = "") :
    loss_arr = []
    acc_arr = []
    for epoch in range(epochs):
        y_hat = model(X_train)
        loss = F.cross_entropy(y_hat , Y_train)
        loss_arr.append(loss.item())
        acc_arr.append(accuracy(y_hat , Y_train))
        loss.backward() #backpropagation
```

```
with torch.no_grad():
    for param in model.parameters():
        param = learning_rate*param.grad
```



```
param -= learning_rate * param.grad
model.zero_grad()

plt.plot(loss_arr , "r-" ,label = "loss")
plt.plot(acc_arr , "b-" ,label = "train_accuracy")
plt.legend()
plt.title(title)
plt.xlabel("Epochs")
plt.show()
print("Loss before training" , loss_arr[0])
print("Loss after training" , loss_arr[-1])

model = FirstNetwork()

fit(10000 , 0.2 , "Loss plot - nn.Parameter & nn.Module")
```

Loss plot - nn Parameter & nn Module

▼ Using Optim and NN linear

```
class FirstNetwork_v1(nn.Module):
    def __init__(self) :
        super().__init__()
        torch.manual_seed(0)
        self.lin1 = nn.Linear(2,2)
        self.lin2 = nn.Linear(2,4)

    def forward(self ,X):
        a1 = self.lin1(X)
        h1 = a1.sigmoid()
        a2 = self.lin2(h1)
        h2 = a2.exp() / a2.exp().sum(-1).unsqueeze(-1)
        return h2

model = FirstNetwork_v1()

fit(10000 , 0.2 , "Loss plot - nn.Linear")
```



```
from torch import optim
```

```
0.4
```

```
def fit_v1(epochs = 10000 , learning_rate = 0.2 , title = "") :
```

```
    loss_arr = []
```

```
    acc_arr = []
```

```
    opt = optim.SGD(model.parameters() , lr = learning_rate)
```

```
    for epoch in range(epochs):
```

```
        y_hat = model(X_train)
```

```
        loss = F.cross_entropy(y_hat , Y_train)
```

```
        loss_arr.append(loss.item())
```

```
        acc_arr.append(accuracy(y_hat , Y_train))
```

```
    loss.backward()
```

```
    opt.step()
```

```
    opt.zero_grad()
```

```
plt.plot(loss_arr , "r-" , label = "loss" )
```

```
plt.plot(acc_arr , "b-" , label = "train_accuracy")
```

```
plt.legend()
```

```
plt.title(title)
```

```
plt.xlabel("Epochs")
```

```
plt.show()
```

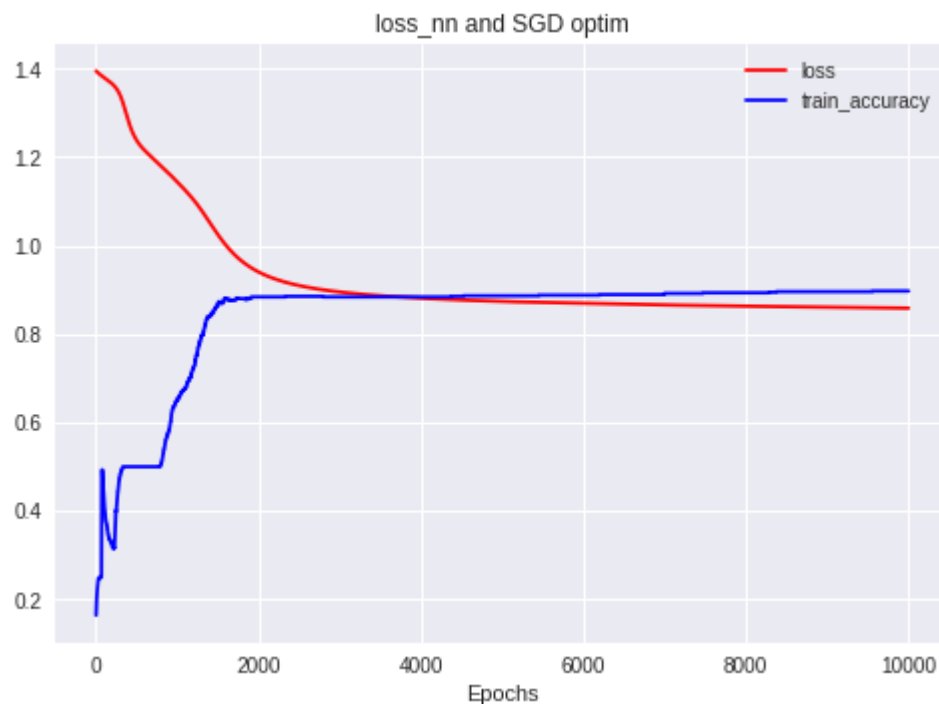
```
print("Loss before training " , loss_arr[0])
```

```
print('Loss after training' , loss_arr[-1])
```

```
%%time
```

```
model = FirstNetwork_v1()
```

```
fit_v1( 10000 , 0.2 , "loss_nn and SGD optim")
```



Loss before training 1.395160436630249

Loss after training 0.8586323857307434

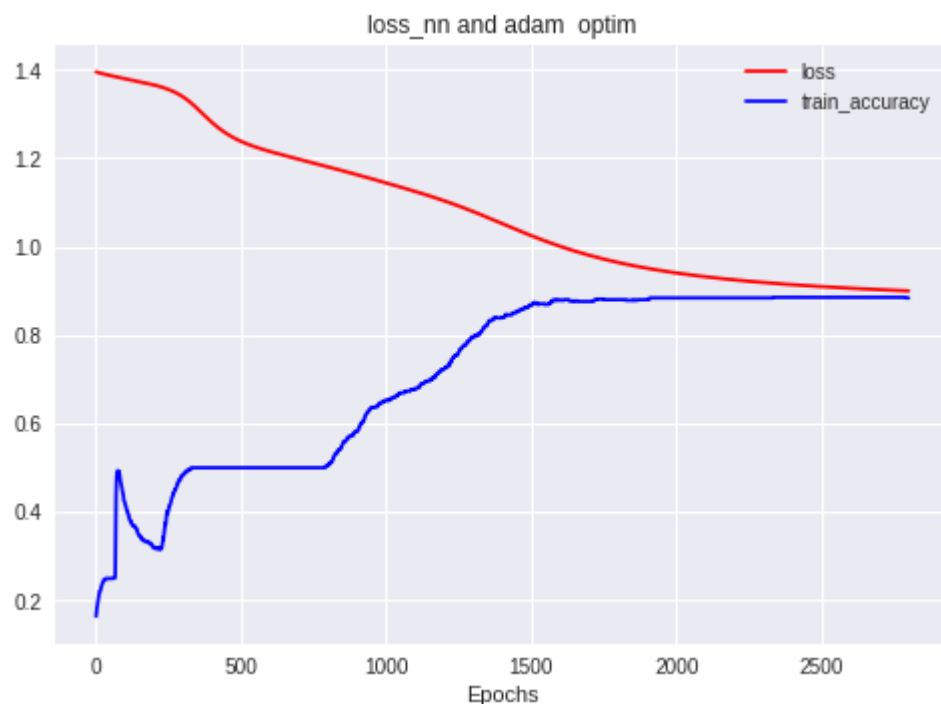
CPU times: user 9.31 s, sys: 48.9 ms, total: 9.36 s

Wall time: 9.36 s

```
%time
```

```
model = FirstNetwork_v1()
```

```
fit_v1( 2800 , 0.2 , "loss_nn and adam optim")
```



Loss before training 1.395160436630249

Loss after training 0.8998849987983704

CPU times: user 2.75 s, sys: 18 ms, total: 2.77 s

Wall time: 2.77 s

▼ Using NN sequential

```
class FirstNetwork_v2(nn.Module):
    def __init__(self):
        super().__init__()
        torch.manual_seed(0)
        self.net = nn.Sequential(
```

```
        nn.Linear(2,2) ,
        nn.Sigmoid() ,
        nn.Linear(2,4) ,
        nn.Softmax() )
def forward(self , X) :
    return self.net(X)


model = FirstNetwork_v2()

def fit_v2(x, y, model, opt, loss_fn, epochs = 10000):
    """Generic function for training a model """
    for epoch in range(epochs):
        loss = loss_fn(model(x), y)

        loss.backward()
        opt.step()
        opt.zero_grad()

    return loss.item()

loss_fn =F.cross_entropy
opt = optim.SGD(model.parameters() , lr = 0.2)

fit_v2(X_train , Y_train , model , opt , loss_fn)
```

0.8586323857307434

▼ Running on GPUs

```
device = torch.device("cuda")
```

```
X_train=X_train.to(device)
```

```
Y_train=Y_train.to(device)
```

```
model = FirstNetwork_v2()
```

```
model.to(device) #moving the network to GPU
```

```
#calculate time
```

```
tic = time.time()
```

```
print('Final loss', fit_v2(X_train, Y_train, model, opt, loss_fn))
```

```
toc = time.time()
```

```
print('Time taken', toc - tic)
```

```
Final loss 1.395159363746643
```

```
Time taken 7.821534633636475
```

```
class FirstNetwork_v3(nn.Module):
```

```
    def __init__(self):
```

```
        super().__init__()
```

```
        torch.manual_seed(0)
```

```
        self.net = nn.Sequential(
```

```
            nn.Linear(2, 1024*4),
```

```
            nn.Sigmoid(),
```

```
            nn.Linear(1024*4, 4),
```

```
            nn.Softmax())
```

```
    def forward(self, X):
```

```
        return self.net(X)
```

```
device = torch.device("cpu")
```

```
X_train=X_train.to(device)
```

```
Y_train=Y_train.to(device)
```

```
#training on gpu
```

```
fn = FirstNetwork_v3()
```

```
fn.to(device)
```

```
tic = time.time()
```

```
print('Final loss', fit_v2(X_train, Y_train, fn, opt, loss_fn))
```

```
toc = time.time()
```

```
print('Time taken', toc - tic)
```

KeyboardInterrupt Traceback (most recent call last)

<ipython-input-34-dee9772d0354> in <module>()

9

10 tic = time.time()

---> 11 print('Final loss', fit_v2(X_train, Y_train, fn, opt, loss_fn))

12 toc = time.time()

13 print('Time taken', toc - tic)

2 frames

/usr/local/lib/python3.6/dist-packages/torch/autograd/__init__.py in backward(tensors, grad_tensors, retain_graph, create_graph, grad_variables)

125 Variable._execution_engine.run_backward(

126 tensors, grad_tensors, retain_graph, create_graph,

--> 127 allow_unreachable=True) # allow_unreachable flag

128

129

KeyboardInterrupt:

SEARCH STACK OVERFLOW

